

Nr. 16

a) Satz von Bayes beweisen:

$$P(S|W) = \frac{P(W|S) \cdot P(S)}{P(W)}$$

bedingte Wahrscheinlichkeit:

$$P(S|W) = \frac{P(S \cap W)}{P(W)} \quad P(W) > 0$$

$$P(W|S) = \frac{P(S \cap W)}{P(S)}$$

$$P(S \cap W) = P(S|W) \cdot P(W) = P(W|S) \cdot P(S)$$

$$\Rightarrow P(S|W) = \frac{P(W|S) \cdot P(S)}{P(W)} \quad \checkmark \quad \square$$

b) Wahrscheinlichkeit, dass heute Fußball gespielt wird:

Attribut	Value	S = yes	S = No
Wind	high	3	3
Humidity	high	3	4
Temperature	cold	3	3
Forecast	Sunny	2	1

$$14 \text{ Messungen} \Rightarrow P(S = \text{yes}) = \frac{9}{14}$$

$$P(S = \text{No}) = \frac{5}{14}$$

} Wahrscheinlichkeit für jede einzelne Wetterbedingung, dass Fußball gespielt wird

$$P(W|S) = \prod_i P(x_i | S)$$

$$P(W|S = \text{yes}) = \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{2}{9} = \frac{2}{243} \quad \checkmark$$

$$P(W|S=No) = \frac{3}{5} \cdot \frac{4}{5} \cdot \frac{3}{5} \cdot \frac{1}{5} = \frac{36}{625} \quad \checkmark$$

Normalizing  $P(W)$ :  $f_{\text{aktuelle Wahrscheinlichkeit}}$

$$P(W) = P(W|S=yes) \cdot P(S=yes) + P(W|S=No) \cdot P(S=No)$$

$$= \frac{2}{243} \cdot \frac{5}{14} + \frac{36}{625} \cdot \frac{5}{14} = \frac{611}{23625} \quad \checkmark$$

$$\begin{aligned} \Rightarrow P(S=yes|W) &= \frac{P(W|S=yes) \cdot P(S=yes)}{P(W)} \\ &= \frac{\frac{2}{243} \cdot \frac{5}{14}}{\frac{611}{23625}} = \frac{125}{611} = 0.2047 = 20.47\% \end{aligned}$$

Mit einer Wahrscheinlichkeit von 20.47% wird heute Fußball gespielt.

c) Wahrscheinlichkeit, dass morgen Fußball gespielt wird

Attribut	Value	S=yes	S=No
Wind	low	6	2
Humidity	high	3	4
Temperature	hot	0	1
Forecast	Sunny	2	1

c)

We read in the data, same as in b).

Notably, there are no games played on a hot day. This results in a probability of `zero`, which raises an issue.

In particular, the value `p_temp` is used to calculate the overall probability in

$$p_W = (p_{W\_yes} * p_S) + (p_{W\_no} * p_{S\_})$$

$$p_{SW} = p_{W\_yes} * p_S / p_W$$

, so it appears in both, the normalization and in the final scaling of `p_SW`.

One possible workaround is the **Add-One Smoothing** or **Laplace Smoothing**. This method adds an arbitrary 1 to each attribute count. (To mitigate those counts which are 0) At the same time, the overall counter is increased by the total number of attributes; to account for the 1-addition.

In formulas, this method can be described as follows:

$$p_{i, \text{empirical}} = \frac{x_i}{N}$$

This is the problematic probability, being zero. We smooth it out by adding

$$p_{i, \alpha\text{-smoothed}} = \frac{x_i + \alpha}{N + \alpha d}$$

as if to increase each count  $x_i$  by  $\alpha$  a priori. (In this case  $\alpha = 1$ ,  $d = 4$ )

Yes, possible and that is what was wanted,  
I am not a fan of this solution.  
Discussion -> see class

```
In [ ]: import pandas as pd

df_c = pd.DataFrame(
    [
        ["Wind", "low", 6, 2],
        ["Humidity", "high", 3, 4],
        ["Temperature", "hot", 0, 1],
        ["Forecast", "sunny", 2, 1]
    ],
    columns=["attribute", "value", "S=yes", "S=no"]
)
df_c.head()

sum_yes = 9 + 4
# Calculate probabilities with Laplace smoothing (instead of omitting values)
p_wind = (df_c.loc[0, "S=yes"] + 1) / sum_yes
p_humidity = (df_c.loc[1, "S=yes"] + 1) / sum_yes
p_temp = (df_c.loc[2, "S=yes"] + 1) / sum_yes
p_cast = (df_c.loc[3, "S=yes"] + 1) / sum_yes

p_W_yes = p_wind * p_humidity * p_temp * p_cast

sum_no = 5 + 4
p_wind_no = (df_c.loc[0, "S=no"] + 1) / sum_no
p_humidity_no = (df_c.loc[1, "S=no"] + 1) / sum_no
p_temp_no = (df_c.loc[2, "S=no"] + 1) / sum_no
p_cast_no = (df_c.loc[3, "S=no"] + 1) / sum_no

p_W_no = p_wind_no * p_humidity_no * p_temp_no * p_cast_no

total = sum_yes + sum_no
p_S = sum_yes / total
p_S_ = sum_no / total

p_W = (p_W_yes * p_S) + (p_W_no * p_S_)

p_SW = p_W_yes * p_S / p_W

print(f"Probability for today: {p_SW:.4f} %")
```

Probability for today: 0.3172 %

As reference: The probability with omitting of `p_temp` is ~58%.

Nr. 17

a) Calculate the entropy of the tree's root

$$H(Y) = - \sum_{z \in Z} P(Y=z) \log_2 P(Y=z) \quad \rightarrow \text{Entropy before a split}$$

- random variable  $Y$

$\rightarrow$  takes values from finite set of symbols  $Z$

$$H(Y|X) = \sum_{m \in M} P(X=m) H(Y|X=m) \quad \rightarrow \text{Entropy after the split}$$

$$= - \sum_{m \in M} P(X=m) \sum_{z \in Z} P(Y=z|X=m) \log_2 P(Y=z|X=m)$$

$$\Rightarrow N_{\text{Hassys}} = 14$$

$$\text{True} = 9$$

$$\text{False} = 5$$

$$H(Y) = H(S) = - \sum_{z \in \{\text{True}, \text{False}\}} P(S=z) \log_2 (P(S=z))$$

$$= - \left[ P(S=\text{True}) \cdot \log_2 (P(S=\text{True})) + P(S=\text{False}) \cdot \log_2 (P(S=\text{False})) \right]$$

$$P(S=\text{True}) = \frac{9}{14}$$

$$P(S=\text{False}) = \frac{5}{14}$$

$$H(S) = - \left[ \frac{9}{14} \cdot \log_2 \left( \frac{9}{14} \right) + \frac{5}{14} \cdot \log_2 \left( \frac{5}{14} \right) \right] = 0,94 \quad \checkmark$$

b) Calculate the information gain if cut is made on the attribute wind

$$I_G(X, Y) = (H(Y) - H(Y|X))$$

wind:  
6 true  
8 false

X = wind  
Y = Soccer

$$H(Y|X) = \sum_{m \in M} P(X=m) H(Y|X=m)$$

$$= \sum_{m \in \{\text{true}, \text{false}\}}$$

$$\begin{aligned} H(S | \text{Wind} = \text{True}) &= - \left[ P(S = \text{true} | \text{Wind} = \text{true}) \cdot \log_2 [P(S = \text{true} | \text{Wind} = \text{true})] \right. \\ &\quad \left. + P(S = \text{false} | \text{Wind} = \text{true}) \cdot \log_2 [P(S = \text{false} | \text{Wind} = \text{true})] \right] \\ &= - \left[ \frac{3}{6} \cdot \log_2 \left( \frac{3}{6} \right) + \frac{3}{6} \cdot \log_2 \left( \frac{3}{6} \right) \right] = 1 \end{aligned}$$

$$H(S | \text{Wind} = \text{false}) = - \left[ \frac{6}{8} \log_2 \left( \frac{6}{8} \right) + \frac{2}{8} \log_2 \left( \frac{2}{8} \right) \right] = 0,811 \checkmark$$

$$\begin{aligned} H(S | \text{Wind}) &= P(\text{Wind} = \text{true}) \cdot H(S | \text{Wind} = \text{true}) + P(\text{Wind} = \text{false}) \cdot H(S | \text{Wind} = \text{false}) \\ &= \frac{6}{14} \cdot 1 + \frac{8}{14} \cdot 0,811 = 0,892 \checkmark \end{aligned}$$

$$I_G(\text{Wind}, S) = H(S) - H(S | \text{Wind})$$

$$= 0,94 - 0,892 = 0,048 \checkmark$$

$$\sum_{A16} = 5/5$$

```

In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("soccer.csv")

d = {'sunny': 0, 'cloudy': 1, 'rainy': 2}
inverse_d = {v: k for k, v in d.items()} # For Later use
df["forecast_code"] = df["weather_forecast"].map(d)

def entropy(target: str, df: pd.DataFrame) -> float:
    """
    The type of the attribute doesn't matter here.
    The cut is done beforehand and this function only handles the subset of the cut.
    The distinction between `soccer==True` and `False` is made in all subsets.
    """
    count_yes = len(df[df[target] == True])
    count_no = len(df[df[target] == False])

    count_sum = count_yes + count_no

    p_yes = count_yes / count_sum
    p_no = count_no / count_sum

    H = - (p_yes * np.log2(p_yes) if p_yes > 0 else 0) - (p_no * np.log2(p_no) if p_no > 0 else 0)

    return H

def information_gain(target: str, df: pd.DataFrame, attribute: str, type: str) -> float:
    """
    Parameters
    -----
    >>> type: ["ordinal", "nominal", "cardinal"]
    ordinal: Ordered, non-numerical values (e.g. low, medium, high).
    nominal: Non-ordered, non-numerical, categorical values (sunny, rainy, cloudy).
    cardinal: Continuous, numerical values (e.g. temperature = 17.8, 21.1, ...).
    """

    # Initialization
    # Total entropy
    total_entropy = entropy(target, df)

    total_samples = len(df)
    weighted_entropy = 0
    best_cut = None

    if (type == "nominal"):
        # The case of nominal values could be done with the One-Hot-Encoding
        # This implementation however iterates over every possible combination of classes.
        # This option is considered viable for a small attribute value count. (in this case 2 and 3)
        unique_values = df[attribute].unique()
        min_weighted_entropy = 1.

        plot_values = []

        for i in unique_values:
            cut_value = i
            left_sub_df = df[df[attribute] == cut_value]
            right_sub_df = df[df[attribute] != cut_value]
            left_sub_entropy = entropy(target, left_sub_df)
            right_sub_entropy = entropy(target, right_sub_df)

            weighted_entropy = (len(left_sub_df) / total_samples * left_sub_entropy +
                               len(right_sub_df) / total_samples * right_sub_entropy)

            if weighted_entropy < min_weighted_entropy:
                min_weighted_entropy = weighted_entropy
                best_cut = cut_value

        plot_values.append((cut_value, total_entropy - weighted_entropy))

    # Create plot
    plt.figure(figsize=(10,5))
    plt.bar([str(val[0]) for val in plot_values], [val[1] for val in plot_values])
    plt.xlabel('Cut')
    plt.ylabel('Information Gain')

```

```

plt.title(f'Information Gain for different cuts of attribute {attribute}')
plt.show()

weighted_entropy = min_weighted_entropy

elif (type == "ordinal"):
    # We don't have this case, so there's no need to implement it.
    print("No ordinal values in data set.")
elif (type == "cardinal"):
    # Do several cuts and determine the lowest weighted_entropy (to maximize the information gain)
    # Determine the cut by taking the mean between two values. For n values, we have n-1 possible cuts.
    df = df.sort_values(by=attribute)
    min_weighted_entropy = 1

    plot_values = []

    for i in range(total_samples - 1):
        if (df.iloc[i].at[attribute] == df.iloc[i+1].at[attribute]):
            continue
        cut_value = (df.iloc[i].at[attribute] + df.iloc[i+1].at[attribute]) / 2
        left_sub_df = df[df[attribute] <= cut_value]
        right_sub_df = df[df[attribute] > cut_value]
        left_sub_entropy = entropy(target, left_sub_df)
        right_sub_entropy = entropy(target, right_sub_df)

        if (len(left_sub_df) == 0 or len(right_sub_df) == 0):
            print(f"Null subset at {df.loc[i].at[attribute]}, i:{i}")

        weighted_entropy = (len(left_sub_df) / total_samples * left_sub_entropy +
                             len(right_sub_df) / total_samples * right_sub_entropy)

        if weighted_entropy < min_weighted_entropy:
            min_weighted_entropy = weighted_entropy
            best_cut = cut_value

    plot_values.append((cut_value, total_entropy - weighted_entropy))

    # Create plot
    plt.figure(figsize=(10,5))
    plt.bar([val[0] for val in plot_values], [val[1] for val in plot_values])
    plt.xlabel('Cut')
    plt.ylabel('Information Gain')
    plt.title(f'Information Gain for different cuts of attribute {attribute}')
    plt.show()

    weighted_entropy = min_weighted_entropy

    return best_cut, total_entropy - weighted_entropy

df = df.sort_values("temperature", ascending=False)
df

```

	temperature	weather_forecast	humidity	wind	soccer	forecast_code
<b>0</b>	29.4	sunny	85	False	False	0
<b>2</b>	28.3	cloudy	78	False	True	1
<b>12</b>	27.2	cloudy	75	False	True	1
<b>1</b>	26.7	sunny	90	True	False	0
<b>9</b>	23.9	rainy	80	False	True	2
<b>10</b>	23.9	sunny	70	True	True	0
<b>7</b>	22.2	sunny	95	False	False	0
<b>11</b>	22.2	cloudy	90	True	True	1
<b>13</b>	21.7	rainy	80	True	False	2
<b>3</b>	21.1	rainy	96	False	True	2
<b>8</b>	20.6	sunny	70	False	True	0
<b>4</b>	20.0	rainy	80	False	True	2
<b>5</b>	18.3	rainy	70	True	False	2
<b>6</b>	17.8	cloudy	65	True	True	1

a)

$$H(Y) = - \sum_{z \in Z} P(Y = z) \log_2 P(Y = z)$$

```
In [ ]: H_root = entropy("soccer", df)
print(f"Entropy root: {H_root}\n")
```

Entropy root: 0.9402859586706311

b)

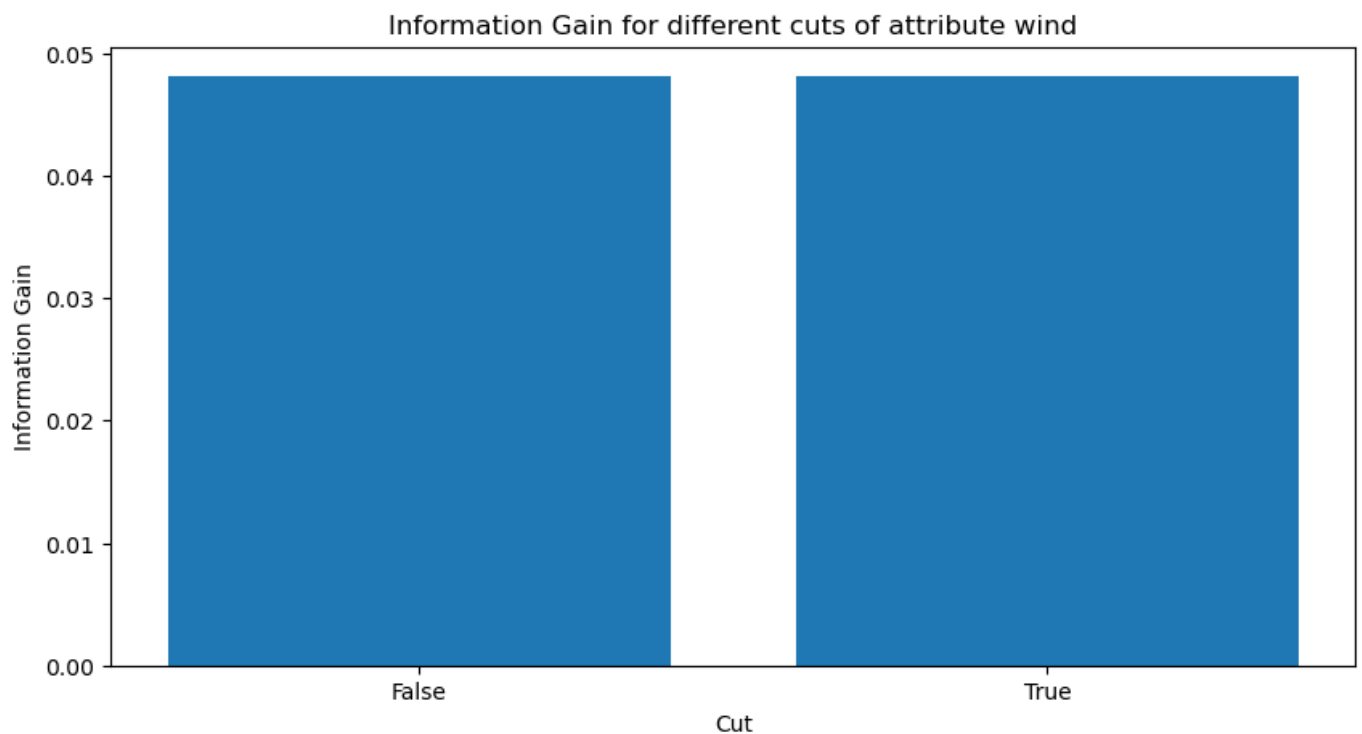
Information gain:  $IG(X, Y) = H(Y) - H(Y | X)$

with

$$\begin{aligned}
 H(Y | X) &= \sum_{m \in M} P(X = m) H(Y | X = m) \\
 &= - \sum_{m \in M} P(X = m) \sum_{z \in Z} P(Y = z | X = m) \log P(Y = z | X = m)
 \end{aligned}$$

```
In [ ]: cut_wind, information_gain_wind = information_gain("soccer", df, "wind", "nominal")
print(f"Information gain wind: {information_gain_wind}\n")
```





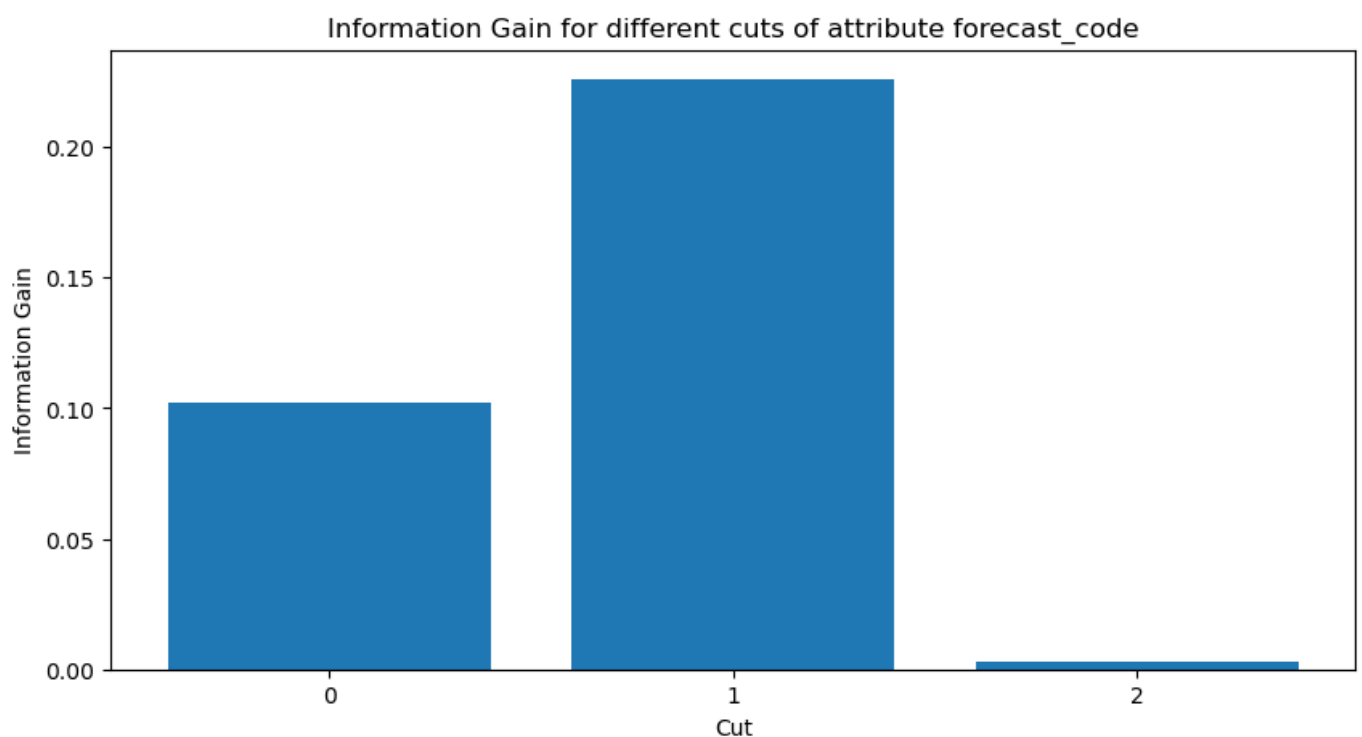
Information gain wind: 0.04812703040826949

c)

```
In [ ]: cut_forecast, information_gain_forecast = information_gain("soccer", df, "forecast_code", "nominal")
print(f"Information gain forecast: {information_gain_forecast}")
print(f"Best cut forecast: {inverse_d[cut_forecast]}\n")


cut_temperature, information_gain_temperature = information_gain("soccer", df, "temperature", "cardinal")
print(f"Information gain temperature: {information_gain_temperature}")
print(f"Best cut temperature: {cut_temperature}\n")

cut_humidity, information_gain_humidity = information_gain("soccer", df, "humidity", "cardinal")
print(f"Information gain humidity: {information_gain_humidity}")
print(f"Best cut humidity: {cut_humidity}\n")
```




Information gain forecast: 0.22600024438491684  
Best cut forecast: cloudy

Information Gain for different cuts of attribute temperature



Information gain temperature: 0.1134008641811034  
Best cut temperature: 28.85

Information Gain for different cuts of attribute humidity



Information gain humidity: 0.10224356360985076  
Best cut humidity: 82.5



d)

## Answer: weather forecast

The best attribute to derive a decision is the **weather forecast**. The information gain is highest when a cut is made at the attribute value "**cloudy**".

A quick glance at the data shows why: On **all** cloudy days, the soccer game was played. The separation of the (sub) data sets between `soccer==True` and `soccer==False` is the "cleanest" (lowest entropy) for this attribute. ✓

$$\sum_{A17} = 5/5$$